



# Redis Data Integration

Hands-on workshop

# Agenda

---

RDI Architecture

---

The data pipeline

---

REST api

---

Installing RDI

---

Creating data pipelines

---

Running RDI

---

# Learning Objectives

By the end of this learning you will be able to...

- **Explain** what Redis Data Integration is and how it works
- **Install** Redis Data Integration in a virtual machine
- **Configure** Redis Data Integration for change data capture (CDC) ingestion
- **Run** and monitor Redis Data Integration

# Workshop Format

This workshop will combine discussion of the principles and practice of RDI with hands-on exercises in a VM that we provide.

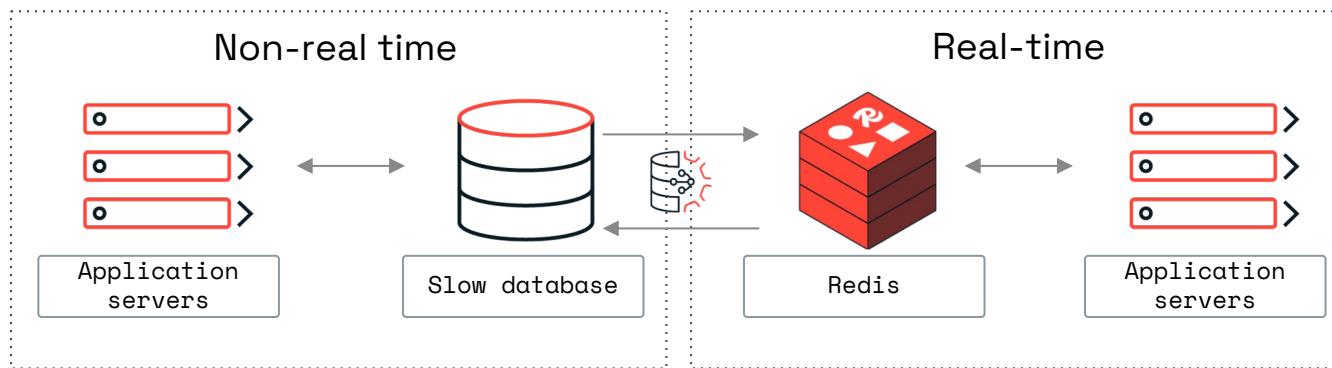
# Redis Data Integration (RDI)

- Redis Data Integration (RDI) is a tool that runs alongside Redis Software
- It lets developers and architects synchronize existing databases with Redis Software to mirror application data
- It helps devs and architects save time on their data integration efforts



# How RDI works

- It creates a data streaming pipeline that mirrors data from an existing database to Redis
- RDI's ingest flow capabilities include a Change Data Capture (CDC) platform to capture changes in the source database
- Within the RDI process, the data is filtered, transformed, and mapped to Redis data types such as Hash or JSON
- RDI then writes the data to the destination Redis database-and other applications can read from the Redis database without complicated workarounds



# When to use RDI

You should use RDI when:

- You must use a slow database as the system of record for the app.
- The app must always *write* its data to the slow database.
- You already intend to use Redis for the app cache.
- The data changes frequently in small increments.
- Your app can tolerate *eventual* consistency of data in the Redis cache.

You should not use RDI when:

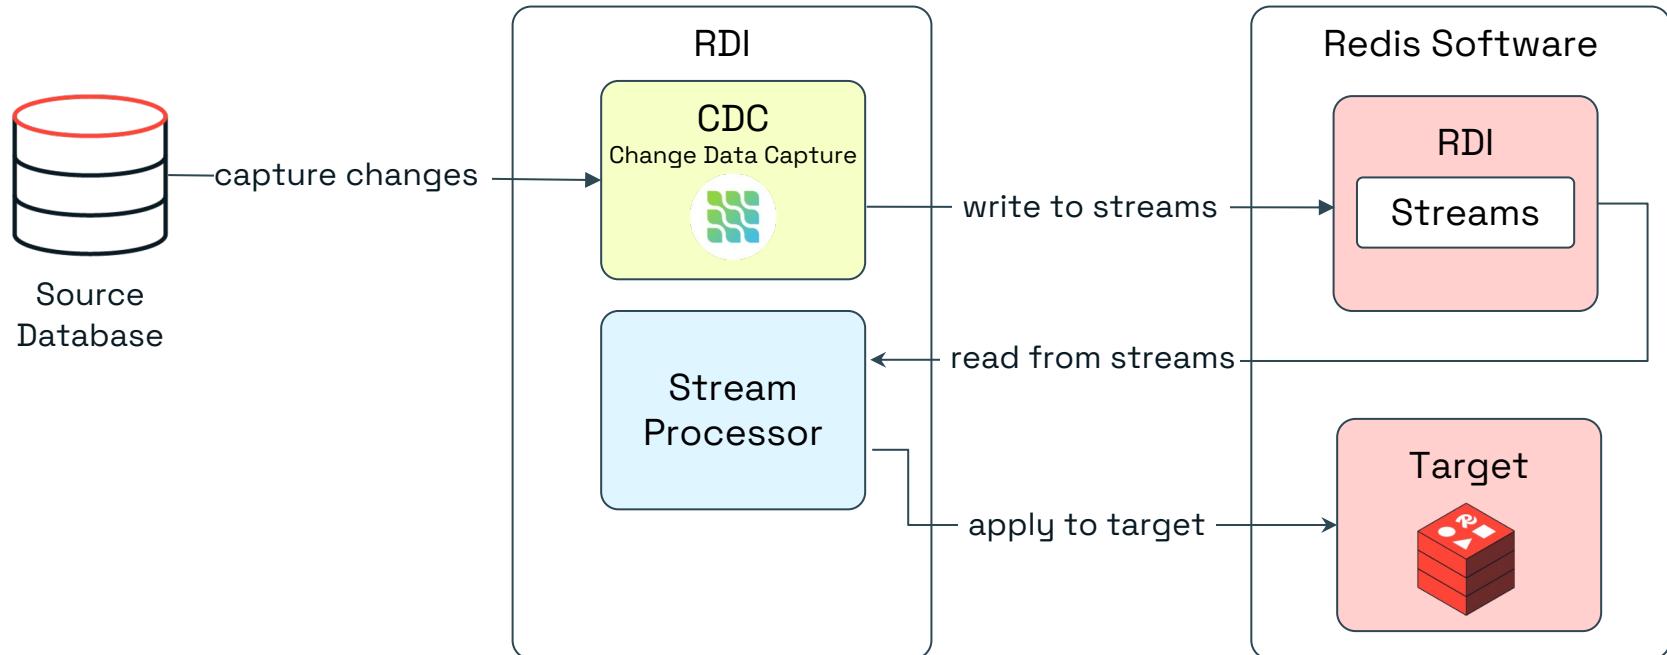
- You are migrating an existing data set into Redis only once.
- The data is updated infrequently and in big batches.
- Your app needs *immediate* cache consistency rather than *eventual* consistency.
- The app must *write* data to the Redis cache, which then updates the source database.
- Your data set will only ever be small.

sd2151kd1fq05m18iq10zmahtrw2mvmlbasj35jk112pldfu  
n457n39dife29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8r  
t9yuip6dfg4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215  
ikd17**thankyou**.rw2mvmlbasj35jk112pldfun457n39dif  
e29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg  
4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m1  
8iq10zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9  
x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk4  
11mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m18iq10  
zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g  
66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb  
2vcx **Redis** d5f68fet4y5h6sd215ikd1fg05m18iq10zmdhf  
rw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g66rga  
7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb2vcx3

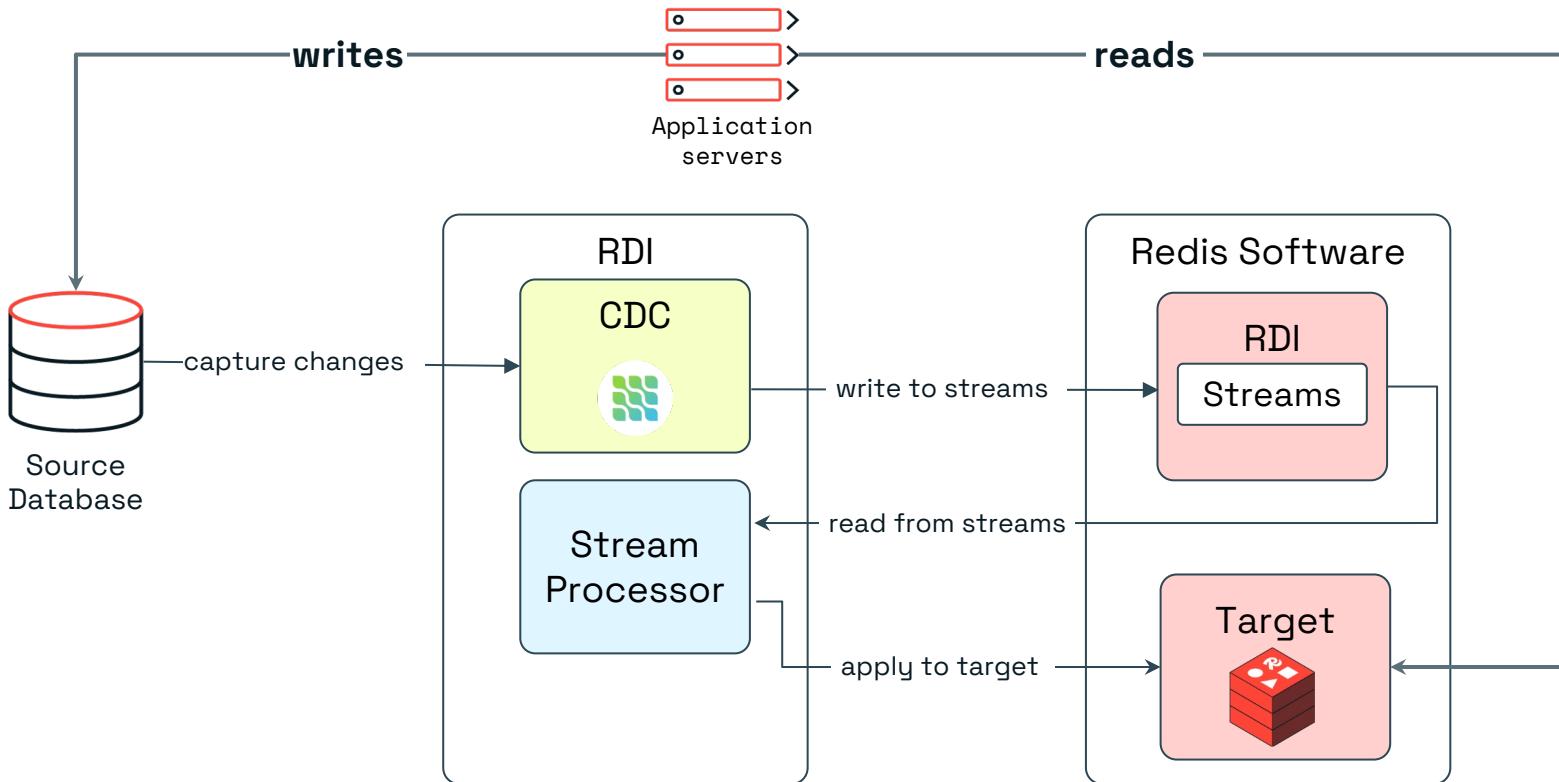
# Architecture Overview

# RDI deployment

## Data plane

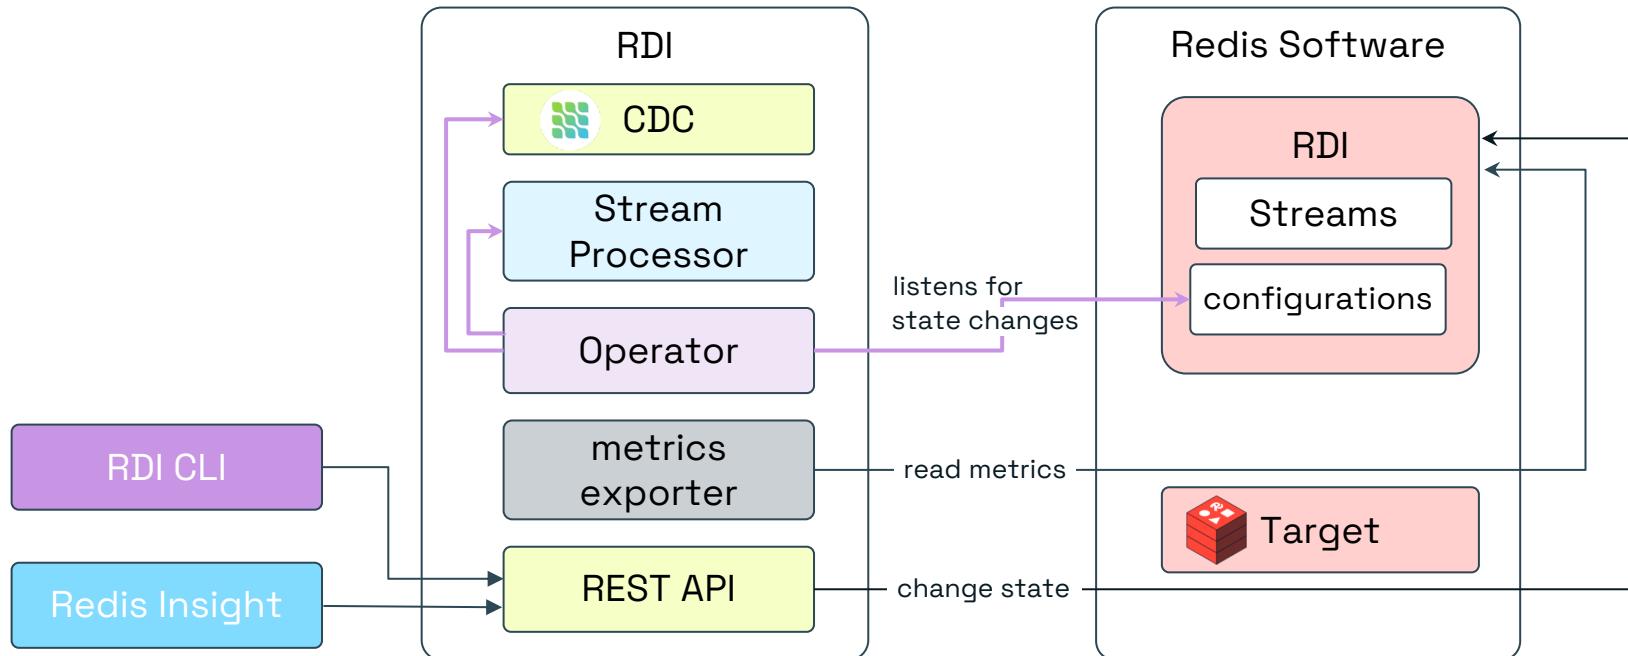


# Redis Data Integration with your apps



# RDI Deployment

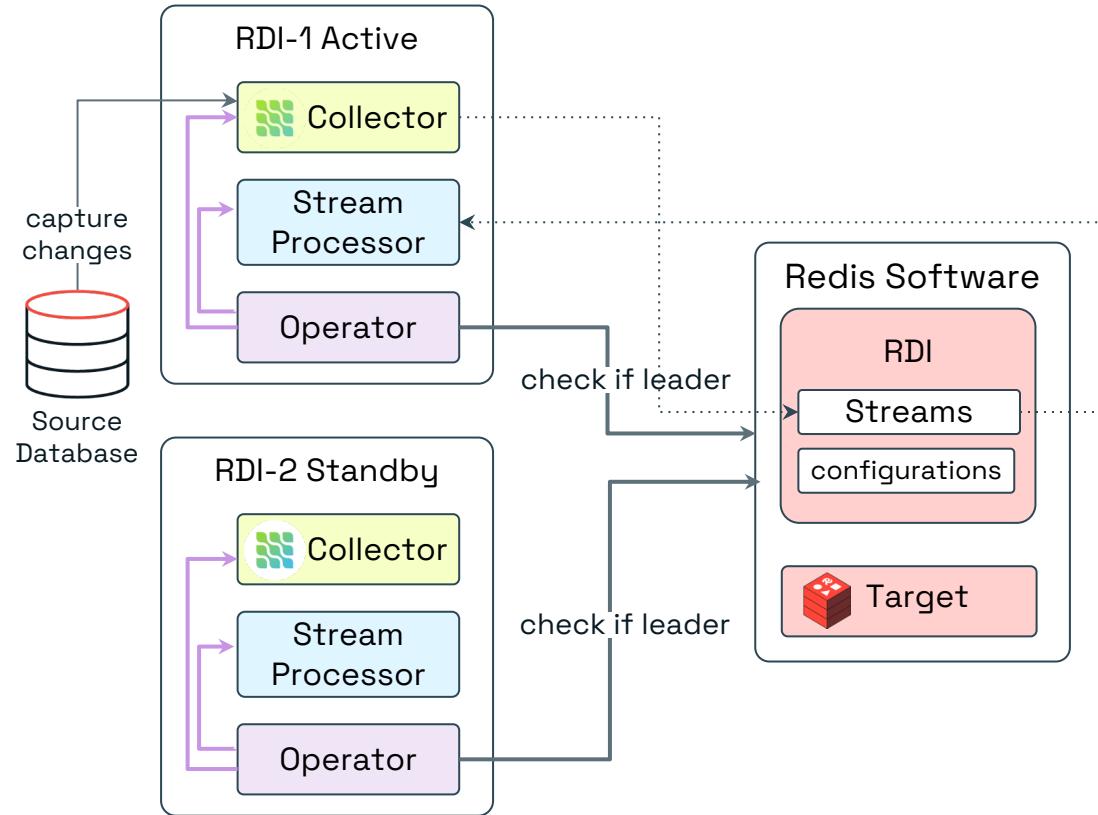
Control plane



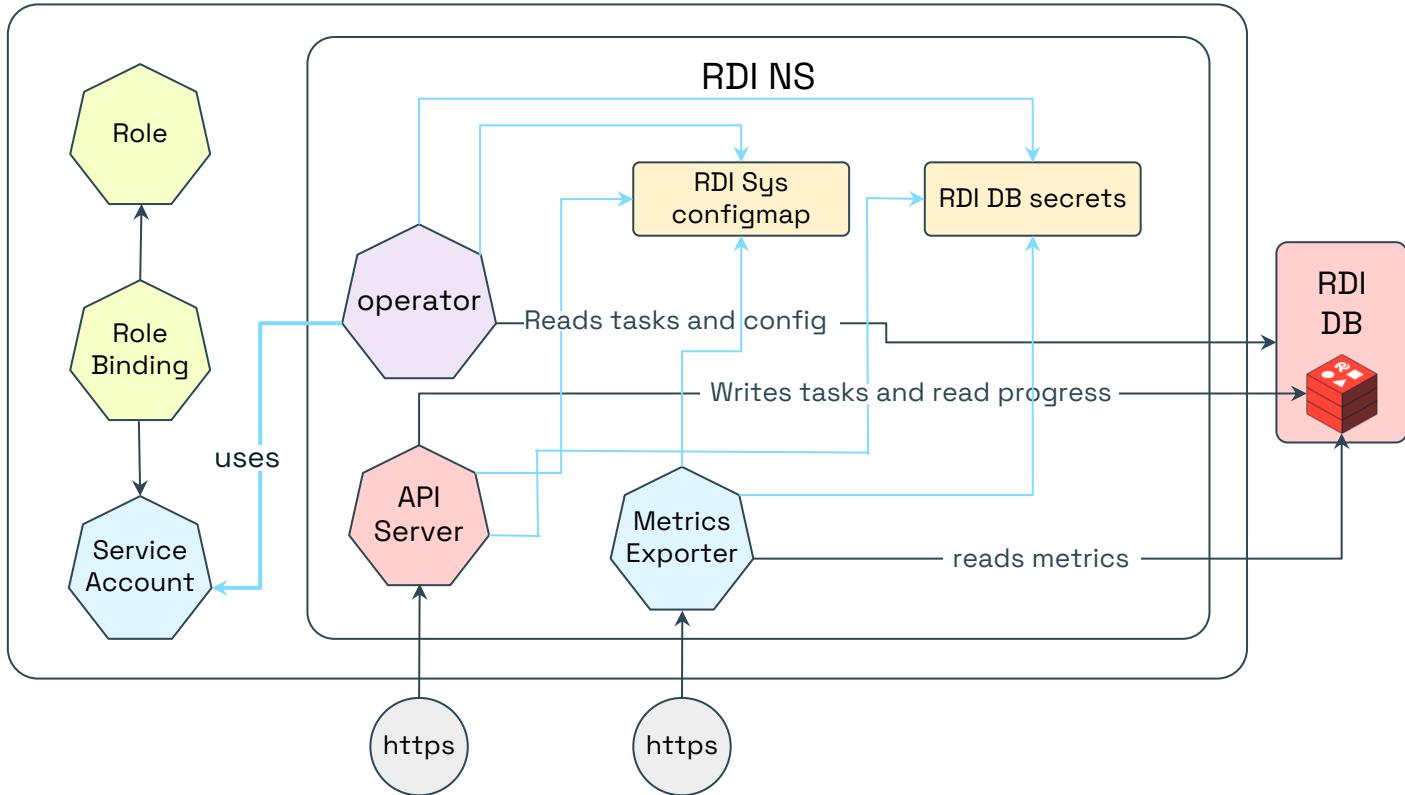
# RDI on VMs

High availability (HA) deployment

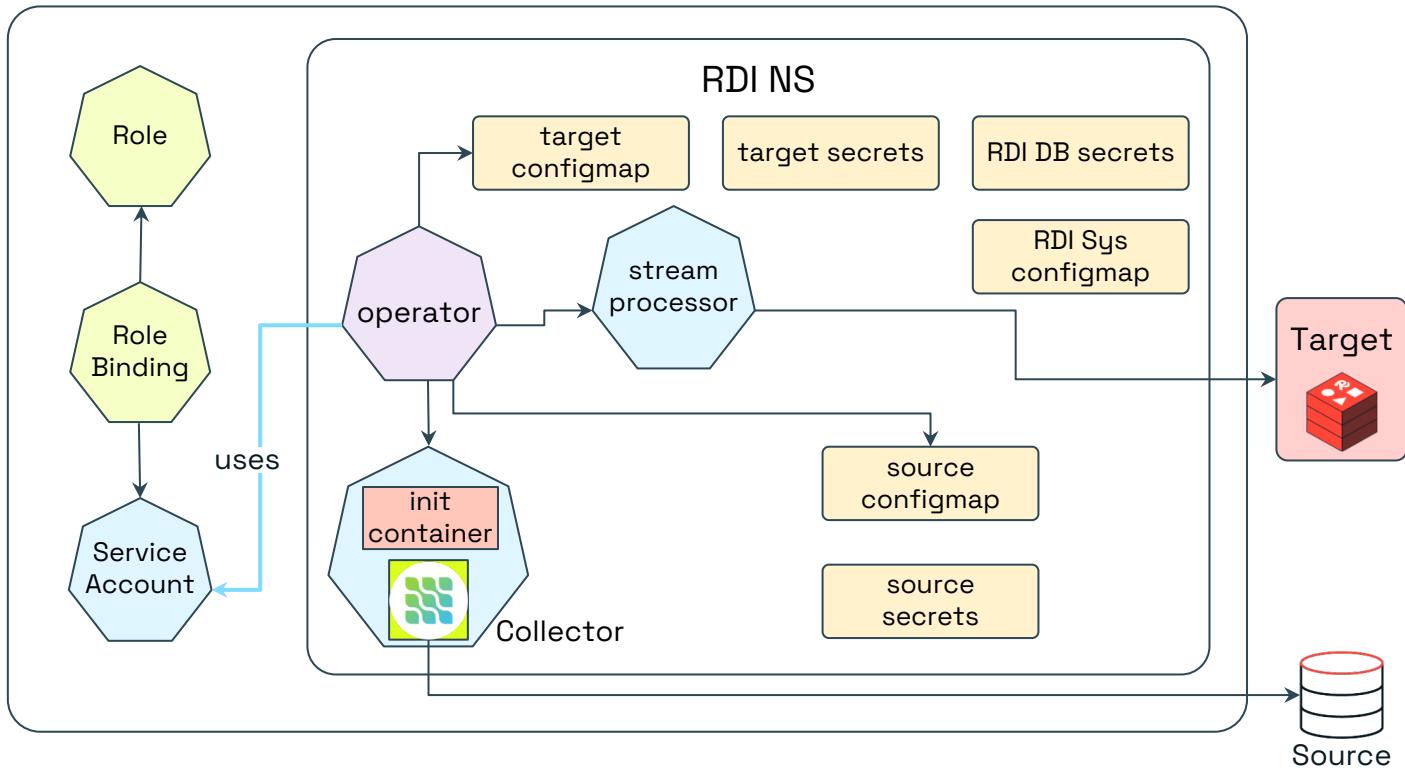
- The operators run on both VMs and use an algorithm to decide which is the active one (the "leader").
- The collector and stream processor are active on one VM and the other is a standby to provide HA.
- You can also install on just one VM if you don't need HA, such as in dev.



# RDI control plane - K8s



# RDI data plane - K8s



sd2151kd1fq05m18iq10zmahtrw2mvmlbasj35jk112pldfu  
n457n39dife29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8r  
t9yuip6dfg4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215  
ikd17**thankyou**.rw2mvmlbasj35jk112pldfun457n39dif  
e29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg  
4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m1  
8iq10zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9  
x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk4  
11mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m18iq10  
zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g  
66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb  
2vcx **Redis** d5f68fet4y5h6sd215ikd1fg05m18iq10zmdhf  
rw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g66rga  
7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb2vcx3

# Prepare Databases

# Before installation

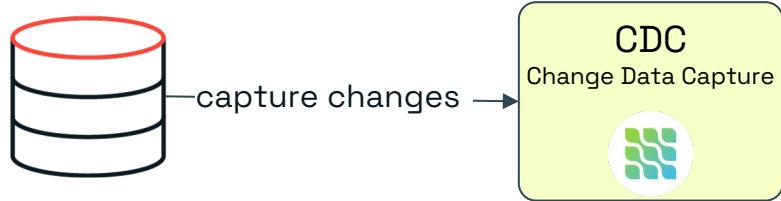
Complete the following steps before you can successfully install/configure RDI:

- Create the RDI database on your Redis Enterprise cluster.
- Create a user for the RDI database if you prefer not to use the default password (see Access control for more information).
- Prepare the source database (e.g. Postgres)

# Preparing the source DB

The first step is to connect CDC to the source DB

- A CDC collector captures changes to the source database.
  - RDI currently uses an open source collector called Debezium for this step.
- The Debezium implementation depends on the source DB. Examples include:
  - Mining the source commit log
  - Querying binlog events
  - Monitoring change publications via plugin
  - Using a stored procedure in the source db



Instructions for supported connections

- [Oracle and Oracle RAC for RDI](#)
- [MySQL/MariaDB for RDI](#)
- [PostgreSQL for RDI](#)
- [SQL Server for RDI](#)

# Preparing Source for the RDI Collector

This is where most problems are created!

- You need to verify the following.
  - Create the debezium DB user with the right permissions for the right db/schema
  - Enable the CDC mechanism (different between db types) for the schema/required tables
  - There is sufficient space for the snapshot and the commit log
- Debezium needs unique row id to generate CDC events.
  - Make sure tables have a PK or unique constraint
  - If one doesn't, you can configure config.yaml for a composite key for this table
    - Use message.key.columns in config.yaml to set the expression

sd2151kd1fq05m18iq10zmahtrw2mvmlbasj35jk112pldfu  
n457n39dife29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8r  
t9yuip6dfg4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215  
ikd17**thankyou**.rw2mvmlbasj35jk112pldfun457n39dif  
e29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg  
4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m1  
8iq10zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9  
x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk4  
11mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m18iq10  
zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g  
66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb  
2vcx **Redis** d5f68fet4y5h6sd215ikd1fg05m18iq10zmdhf  
rw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g66rga  
7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb2vcx3

# Installing RDI

# Before you start...

## Important information

- Download the appropriate RDI tar file from the download center.
- You can't install RDI on a host where a Redis Enterprise cluster is also installed, due to incompatible network rules.
- The supported OS versions for RDI are
  - RHEL 8 & 9
  - Ubuntu 18.04 & 20.04
- Helm charts are supported
- For VMs, you must run the RDI installer as a privileged user because it installs containerd and registers services. (You don't need any special privileges to run RDI processes for normal operation.)
- Additional requirements and considerations can be found in the documentation
  - [Doc for installing on VMs](#)
  - [Doc for installing on Kubernetes](#)

# Installation steps

Follow the steps below for each of your VMs:

- Download the RDI installer from the Redis download center (under the “Modules, Tools & Integration” dropdown)
- Go to the installation folder:  
`cd rdi_install/$RDI_VERSION`
- Run the installer as a privileged user:  
`sudo ./install.sh`
- The installer gives you instructions to help you create secrets and create your pipeline.
- You can use the `-f` option to supply answers to the installer’s questions automatically using properties from a TOML file:  
`./install.sh -f silent.toml`

# Using TLS or mTLS

Sample file snippet for silent.toml

```
[rdi.database]
host = "localhost"
port = 12001
username = "username"
password = "password"
ssl = true
```

*continued in next column*

```
# if the RDI database uses TLS/mTLS.
[rdi.database.certificates]
ca = "/home/ubuntu/rdi/certs/ca.crt"
cert = "/home/ubuntu/rdi/certs/client.crt"
key = "/home/ubuntu/rdi/certs/client.key"
passphrase = "foobar"
```

*SSL (true/false)? - If ssl = false then RDI will ignore the settings in the rdi.database.certificates section.*

# K8s Installation

Use the RDI Helm chart to install on Kubernetes

The installation creates the following K8s objects:

- A K8s namespace named rdi.
- Deployments for the RDI operator, metrics exporter, and API server.
- A service account along with a role and role binding for the RDI operator.
- A Configmap for the different components with RDI Redis database details.
- Secrets with the RDI Redis database credentials and TLS certificates.

# Installation steps

## Install the RDI Helm chart

- Decompress the tar file:  
`tar -xvf rdi-<rdi-tag>.tar.gz`

- Open the `values.yaml` file and set the appropriate values for your installation.
  - See the [RDI Kubernetes doc](#) for details.

- Start the installation:

```
helm install <The logical chart name> ./rdi --create-namespace -n rdi
```

- To check the installation, run the following command:

```
helm list -n monitoring -n rdi
```

- You can verify that the RDI API works by adding the server in Redis Insight.

# Installation Troubleshooting

## Common pitfalls

- Install didn't run as sudo (VM installation)
- Problem accessing the RS cluster (credentials, network, DNS)
- Not enough disk space on VM
  - RDI needs several GBs on top of the OS.
  - Ubuntu is recommended to have a minimal disk space of 25GB.

# Installation Troubleshooting

## How to troubleshoot

- RDI should provide a clear indication what went wrong
- Try network commands, such as these, to test connectivity and see what's happening in the network.
  - telnet
  - nslookup
  - traceroute
- Check disk availability
- Troubleshooting RDI components:
  - Look at `/opt/rdi/logs` to see the components' logs in case one or more are not ready
  - Run `kubectl -n rdi get all` to ensure all deployments have their pods ready

# Uninstall RDI

If you want to remove your RDI installation, go to the installation folder and run the uninstall script as a privileged user:

```
sudo ./uninstall.sh
```

The script will verify you are sure before you proceed:

```
This will uninstall RDI and its dependencies, are you sure? [y, N]
```

If you type anything other than "y" here, the script will abort without making any changes to RDI or your source database.

sd2151kd1fq05m18iq10zmahtrw2mvmlbasj35jk112pldfu  
n457n39dife29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8r  
t9yuip6dfg4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215  
ikd17**thankyou**.rw2mvmlbasj35jk112pldfun457n39dif  
e29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg  
4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m1  
8iq10zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9  
x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk4  
11mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m18iq10  
zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g  
66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb  
2vcx **Redis** d5f68fet4y5h6sd215ikd1fg05m18iq10zmdhf  
rw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g66rga  
7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb2vcx3

Data flow

## Data pipeline

# Ingest Overview

## Capture & Stream Changes

- CDC for most popular databases out of the box
- Can integrate with any CDC product / Streaming data source
- Modes of work:
  - Initial baseline snapshot
  - Streaming changes
- At least once guarantee

## Transform & Denormalize

- Automatically converts from source types to Redis strings
- Declarative transformations can be applied per table
- Handles Many to one relationships as JSON
- Handles One to one relationships as Hash or JSON

## Handle Delivery

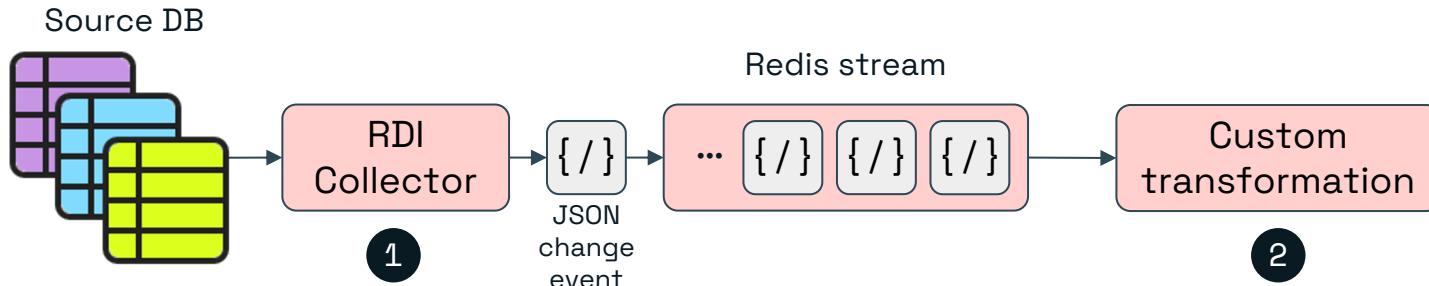
- Supports one source entry to many Redis keys
- Supports Hash, Json, Set, Sorted Set, Stream
- At least once guaranteed
- Malformed data saved to DLQ
- Auto recovers disconnects
- Data provenance

Data flow

Capture and stream data

# Capture & stream changes

## Steps



- A. Initial snapshot (*everything is a change*)
- B. Streaming changes

At least once delivery guarantee

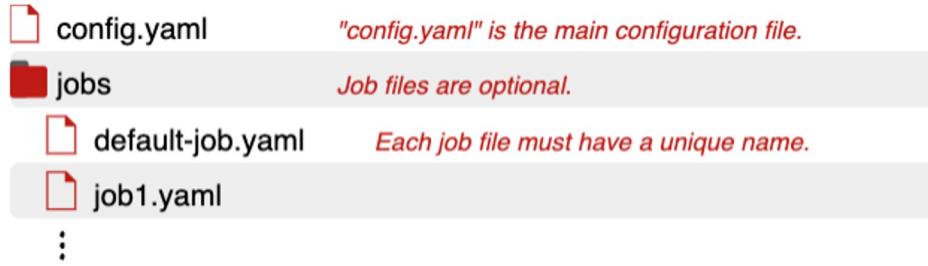
The data transformation involves two separate stages.

1. First, the data ingested during CDC is automatically transformed to a JSON format.
2. Then, this JSON data gets passed on to your custom transformation for further processing.

# Pipeline configuration

Specify the full name of the source table in the job file (or '\*' for default-job.yaml)

RDI uses a set of YAML files to configure each pipeline. The following diagram shows the folder structure:



- The main configuration for the pipeline is in the `config.yaml` file. This specifies:
  - Connection details for the source database
  - Queries that RDI will use to extract the required data.
- You should place job configurations in the `jobs` folder to specify your own data transformations.
- `default-job.yaml` will be used for any table that doesn't have its own `.yaml`

# config.yaml

## Part 1: sources

### The sources section

- A unique name for each source
- type: currently cdc is the only option
- Logging level is up to you
- connection: supply the necessary info to connect to the source
- Include db credentials and TLS/mTLS secrets if applicable

```
sources:  
  my_name:  
    type: cdc  
    logging:  
      level: info  
    connection:  
      type: mysql  
      host: ${RDI_REDIS_HOST}  
      port: 13000  
      database: redislabscdc  
      user: ${SOURCE_DB_USERNAME}  
      password: ${SOURCE_DB_PASSWORD}
```

# config.yaml

## Part 2: sources continued

### The sources section for TLS/mTLS

- The names of the following properties should match the ones you used when setting the TLS/mTLS secrets.
- Set only `cacert` if you are using TLS.
- Set all of them if you are using mTLS.

connection:

*(continuation from previous slide)*

```
key: ${SOURCE_DB_KEY}  
cert: ${SOURCE_DB_CERT}  
cacert: ${SOURCE_DB_CACERT}  
key_password: ${SOURCE_DB_KEY_PASSWORD}
```

# config.yaml

## Part 3: sources continued

### The sources section - tables

- The name of the this table query
- **snapshot\_sql**: A query that selects the table(s) to include in the dataset
  - The default is **all** tables in the db
- **columns**: which ones you want to include in the target
  - The default is **all** columns
- **keys**: A list of primary keys, one for each table
  - or a unique composite if no column has the PRIMARY KEY or a UNIQUE constraint

```
tables:  
  emp:  
    snapshot_sql: "SELECT * from  
redislabscdc.emp WHERE empno < 1000"  
    # the preceding is all one line  
    columns:  
      - empno  
      - fname  
      - lname  
    keys:  
      - empno
```

# config.yaml

## Part 4: target

### The targets section

- A unique name for each target
- connection: info for the target
  - type: presumably redis
  - host: and port: as usual
  - connection credentials and TLS/mTLS secrets if applicable

```
targets:  
  my-redis:  
    connection:  
      type: redis  
      host: localhost  
      port: 12000  
      user: ${TARGET_DB_USERNAME}  
      password: ${TARGET_DB_PASSWORD}
```

# config.yaml

## Part 5: target continued

### The target section for TLS/mTLS

- The names of the following properties should match the ones you used when setting the TLS/mTLS secrets.
- Set only ‘cacert’ if you are using TLS.
- Set all of them if you are using mTLS.

targets:

*(continuation from previous slide)*

key: \${TARGET\_DB\_KEY}

cert: \${TARGET\_DB\_CERT}

cacert: \${TARGET\_DB\_CACERT}

key\_password: \${TARGET\_DB\_KEY\_PASSWORD}

# Transform & denormalize

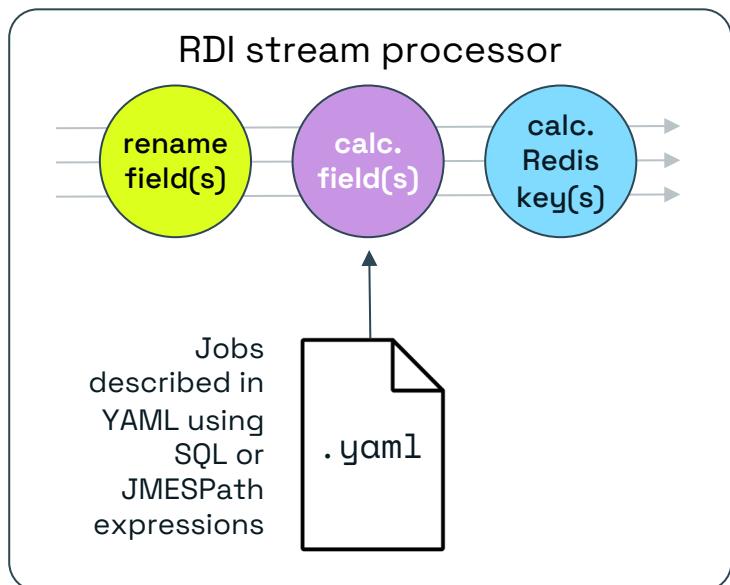
## Single table transformation

Pipeline captures change data records from the source db, and transforms them into Redis data structures. It writes each of these new structures to a Redis target database under its own key.

You can provide your own custom transformation jobs for each source table, using your own data mapping and key pattern.

As part of the transformation, you can specify whether you want to store the data in Redis as JSON objects, hashes, sets, streams, sorted sets, or strings.

## Transformation steps



# Transformation

3 main sections of a job.yaml

## Sample

- source: *(required section)*
- table: Database table name.
  - One of the table names you supplied in config.yaml.
  - Use "\*" for the default job only.
- row\_format: Format of the data to be transformed. Options are:
  - data\_only *(default)*
  - full *(See [the doc](#) for details of the extra data you can access when you use the full option.)*

## Explanation

```
source:  
table: "*"  
row_format: full  
transform:  
    Discussed in next slides  
output:  
    Discussed after transform
```

# Transformation

How to transform data from source to output

## Sample

```
transform: (optional section)
- uses: (see options in next column)
  with: (options depend on use)
```

## Explanation

- **transform:** *(optional section)*
- - **uses:** *(note the dash)*
  - add\_field
  - remove\_field
  - rename\_field
  - filter
  - map
- **with:** *same indentation as uses, not as the dash*

# Adding a field to the output

You can add one or many fields

## Sample

```
transform:  
  - uses: add_field  
    with:  
      fields:  
        - field: name.first_name  
          language: jmespath  
          expression: concat( ... )  
        - field: country  
          language: sql  
          expression: c_code || ...
```

## Explanation

- - uses: add\_field
- with: *(required)*
- fields: *skip if only adding 1 field*
- - field: name of added field in the correct format for the target data type *(skip dash if only 1 field)*
- language: of the expression. Options are jmespath and sql.
- expression: that specifies the value of the new field

# Removing a field from the output

You can remove one or many fields

## Sample

```
transform:  
  - uses: remove_field  
    with:  
      fields:  
        - field: credit_card  
        - field: name.middlename
```

## Explanation

- - uses: remove\_field
- with: *(required)*
- fields: *skip if only removing 1 field*
- - field: name of field to remove in the correct format for the target data type *(skip dash if only 1 field)*

# Rename a field from the output

All other fields remain unchanged

## Sample

```
transform:  
  - uses: rename_field  
    with:  
      fields:  
        - from_field: name.lname  
          to_field: name.last_name  
        - from_field: name.fname  
          to_field: name.first_name
```

## Explanation

- - uses: rename\_field
- with: *(required)*
- fields: *skip if only changing 1 field*
- - from\_field: default name of field  
*(skip dash if only 1 field)*
- to\_field: new name of field *(aligned with from\_field; no dash ever)*

# Map a row to a Redis data type

Map a record into a new output based on expressions

```
transform:  
  - uses: map  
    with:  
      expression: |  
        {  
          "FirstName": first_name,  
          "LastName": last_name,  
          "Location":  
            {  
              "Street": address,  
              "City": city,  
              ... etc. ...  
            }  
        }  
    language: jmespath
```

## Explanation

- - uses: map
- with: *(required)*
- language: of the expression. Options are jmespath and sql.
- expression: that specifies the name of the target field and the column for its value

# Filter records on chosen criteria

Only cache relevant information

## Sample

```
transform:  
  - uses: filter  
    with:  
      language: sql  
      from_field: name.lname  
      expression: age>20
```

## Explanation

- - uses: filter
- with: *(required)*
- language: of the expression. Options are jmespath and sql.
- expression: that specifies the criteria for which rows to include

# Output location, type, and keyname

- output: *(required)*
  - uses: redis.write *(required)*
  - with: *(required)*
- connection: *(optional)* name of target defined in config.yaml
- data\_type: target structure *(optional)* choices are:
  - hash, json, set, sorted\_set, stream, and string
- key: *(optional)* override default key naming for target data
- expression: how to form key
- language: [sql | jmespath]

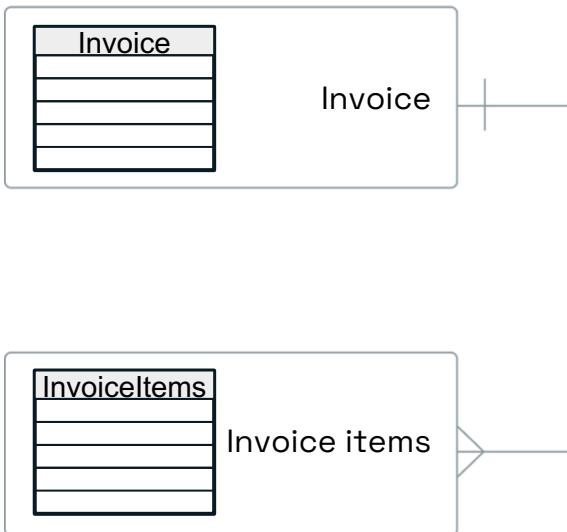
## Sample

```
output:
  - uses: redis.write
    with:
      connection: target
      key:
        expression:
          concat(['emp:fname:', fname, ':lname:', lname])
          # preceding "expression:" is all one line
          language: jmespath
```

# Denormalization

Multi-table transformation with denormalization

Source relational  
database



Many-to-one relationship is  
translated into a single parent  
JSON with a map of child objects

\*..1



```
Invoice:InvoiceID:233 →  
{  
...  
- InvoiceLineItems{  
  1244: {}  
  1156: {}  
...  
}
```

# Sample job for nesting

When normal isn't an option

When you use nesting:

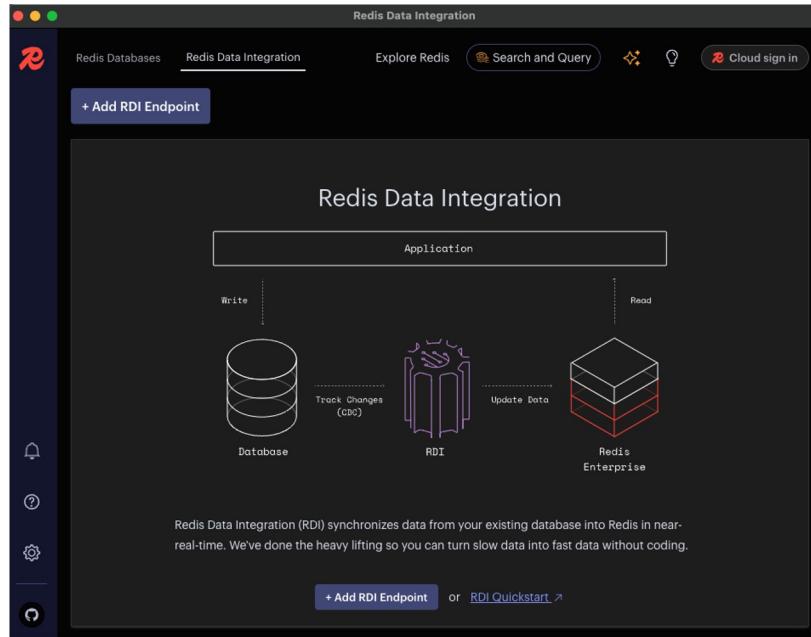
- You must also set the `data_type` attribute to `json` and the `on_update` attribute to `merge` in the output block.
- Key expressions are **not** supported for the nest output blocks.
- You can only use one level of nesting.

```
source:  
  table: InvoiceLine  
output:  
  - uses: redis.write  
    with:  
      nest: # cannot co-exist with other parameters  
      parent:  
        table: Invoice  
      nesting_key: InvoiceLineId # can't be composite  
      parent_key: InvoiceId # cannot be composite  
      path: $.InvoiceLineItems # must start from root  
      structure: map # only map supported for now  
      on_update: merge # only merge supported for now  
      data_type: json # only json supported for now
```

# Deploy a pipeline

There are two options for deploying a pipeline

## Redis Insight



## redis-di CLI

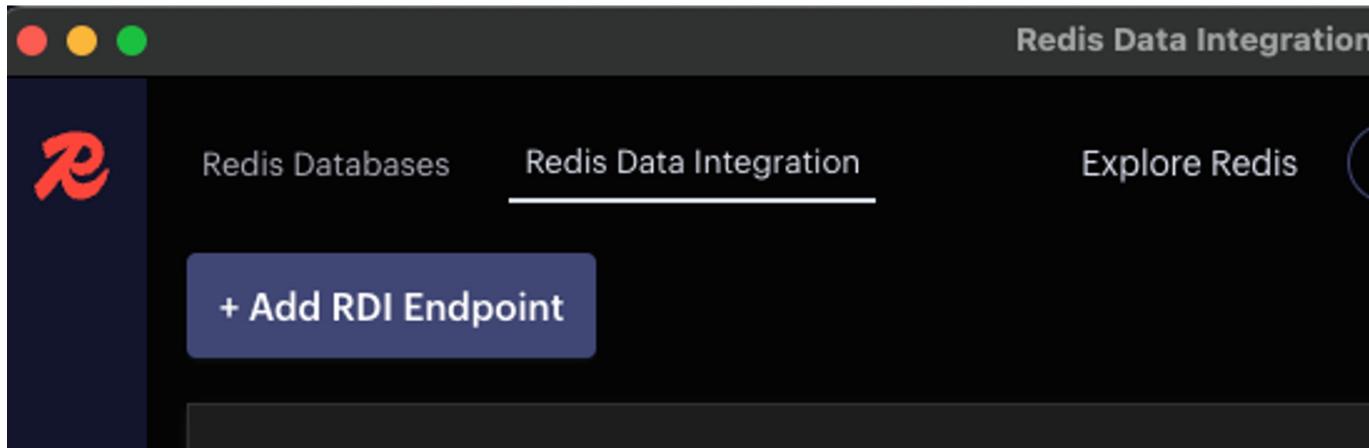
```
redis-di set-secret SOURCE_DB_USERNAME  
myUserName  
# set remaining secrets as needed  
redis-di deploy --dir <path to pipeline folder>
```

# Deploying with Redis Insight

First connect to the RDI server

While on Redis Data Integration page

Click **+Add RDI Endpoint**



# Connect to RDI

Fill in the relevant values

RDI Alias: whatever you want

URL: of the RDI server, not the target db and  
not the source db

Username: to the RDI database that you made

Password: for the RDI database that you made

Click the **Add Endpoint** button

Connect to RDI

RDI Alias\*

URL\*

 ⓘ

Username\*

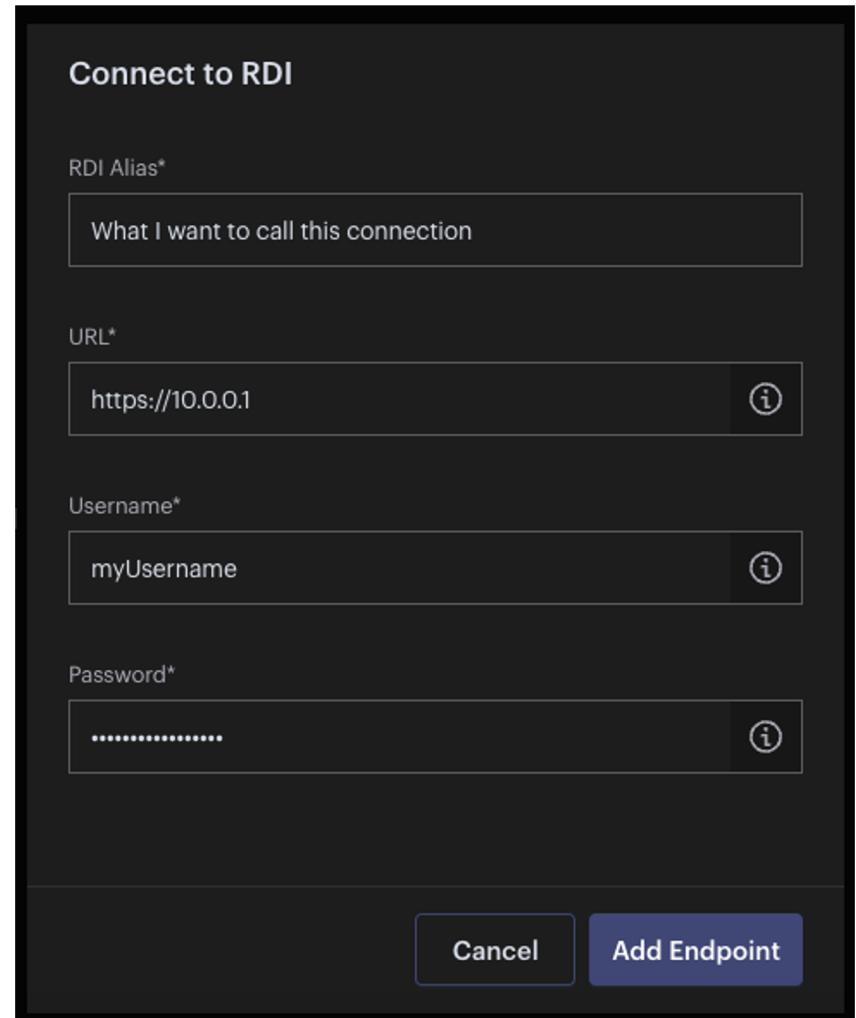
 ⓘ

Password\*

 ⓘ

Cancel

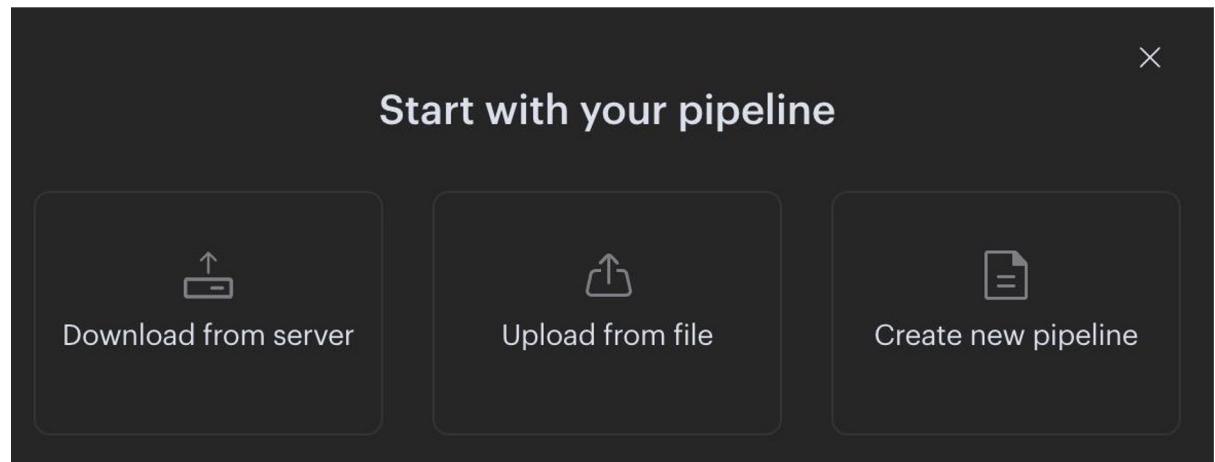
Add Endpoint



# Add the pipeline

You can

- Download from server
- Upload from file
- Start from scratch



This is where you add config.yaml when using Redis Insight.

# Pipeline management

Sample config.yaml

Enter the target and source database information

This shows Redis Insight, but any editor can be used as long as the directory structure is correct.

The screenshot shows the Redis Insight Pipeline Management interface. At the top, it says "RDI instances > Demo". Below that is the "Pipeline Management" section. On the left, there's a sidebar with "Pipeline Management" at the top, followed by "Configure pipeline" and "Configuration file". Under "Add transformation jobs", it shows "Jobs (1)" with a sub-item "track". To the right of the sidebar is a large panel titled "Target database configuration". It contains the text "Configure target instance [connection details](#) and ap" and a code block for config.yaml:

```
targets:
  target:
    connection:
      type: redis
      host: 172.16.22.21
      port: 12000
sources:
  psql:
    type: cdc
    logging:
      level: info
    connection:
      type: postgresql
      host: 172.16.22.7
      port: 5432
      database: chinook
      user: postgres
      password: postgres
```

# Redis Insight lets you test the connection immediately

Target database configuration

Configure target instance [connection details](#) and applier settings.

```
1 targets:
2   .target:
3     .connection:
4       .type: redis
5       .host: 172.16.22.21
6       .port: 12000
7 sources:
8   psql:
9     .type: cdc
10    .logging:
11      .level: info
12      .connection:
13        .type: postgresql
14        .host: 172.16.22.7
15        .port: 5432
```

**Test Connection**

Connection test results

Connected successfully:

1 ✓

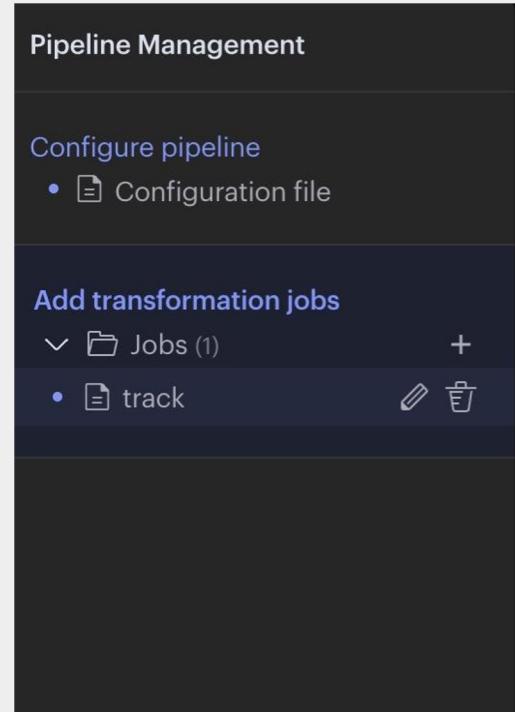
# Transformation jobs

Add as many transformation files as needed for your implementation.

None are required.

By default, RDI transforms the source data into hashes or JSON objects for the target with a standard data mapping and a standard format for the key.

This shows Redis Insight, but any editor can be used as long as the directory structure is correct.

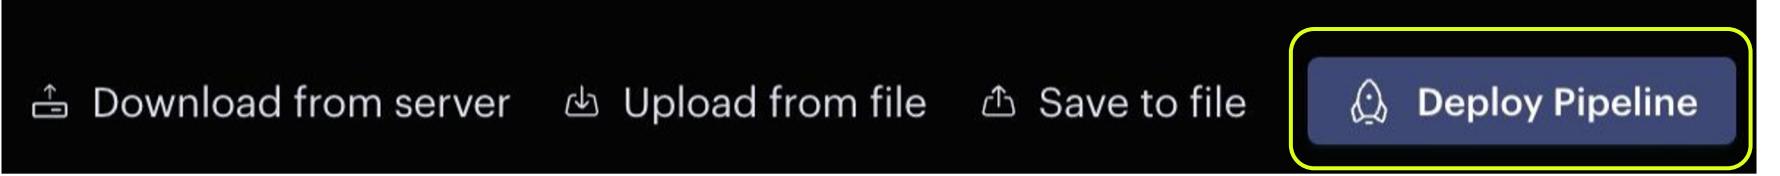


The screenshot shows the Redis Insight Pipeline Management interface. On the left, there's a sidebar with 'Pipeline Management' at the top, followed by 'Configure pipeline' and 'Add transformation jobs'. Under 'Add transformation jobs', there's a section titled 'Jobs (1)' with a plus sign to add more. A single job is listed: 'track'. To the right of the interface, there's explanatory text and a code snippet. The text says 'track' and 'Describe the [transformation logic](#) to perform'. Below this is a multi-line code listing:

```
1  source:
2  |  table: Track
3  transform:
4  |  - uses: add_field
5  |  |  with:
6  |  |  |  field: NameUpper
7  |  |  |  expression: upper("Name")
8  |  |  |  language: sql
9  |  - uses: filter
10 |  |  with:
11 |  |  |  expression: GenreId=2
12 |  |  |  language: sql
13
```

# Deploy the pipeline

Using Redis Insight:



Download from server    Upload from file    Save to file    Deploy Pipeline

A screenshot of the Redis Insight interface. At the bottom, there is a dark bar with three buttons: "Download from server" (with a downward arrow icon), "Upload from file" (with an upward arrow icon), "Save to file" (with an upward arrow icon), and a blue button labeled "Deploy Pipeline" with a rocket ship icon. A yellow rounded rectangle highlights the "Deploy Pipeline" button.

Using CLI:

```
redis-di deploy --dir <path to pipeline folder>
```

sd2151kd1fq05m18iq10zmahtrw2mvmlbasj35jk112pldfu  
n457n39dife29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8r  
t9yuip6dfg4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215  
ikd17**thankyou**.rw2mvmlbasj35jk112pldfun457n39dif  
e29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg  
4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m1  
8iq10zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9  
x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk4  
11mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m18iq10  
zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g  
66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb  
2vcx **Redis** d5f68fet4y5h6sd215ikd1fg05m18iq10zmdhf  
rw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g66rga  
7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb2vcx3

# Observability

# Monitoring RDI

There are three general categories

- Metrics
  - RDI exposes two endpoints, one for CDC collector metrics and another for stream processor metrics.
- Logs
  - By default, RDI stores logs in the host VM file system at /opt/rdi/logs. The logs are recorded at the minimum INFO level and get rotated when they reach a size of 100MB.
- Support package “dump”
  - For a comprehensive set of forensics data you can run `redis-di dump-support-package` from the CLI.

# Pipeline status

Redis Insight has a built in RDI dashboard

The screenshot shows the Redis Insight RDI Lab dashboard. At the top, there's a header with the Redis logo, the title 'RDI Instances > RDI Lab', and various status indicators: RDI Version 7.2.4, Address 172.16.22.21:2001, Run status ready, Sync Mode streaming, and an auto-refresh interval of 5.0 s. A red box highlights the first icon in the top navigation bar, which is a chart symbol.

The main content area is divided into several sections:

- Processing performance information:** Displays metrics like Total batches (211), ACK time average (0.4 ms/sec), Batch size average (1.7 MB), Records per second average (16 /sec), Process time average (5.4 ms), and Read time average (100.2 ms).
- Target connections:** Shows a single connection named 'target' of type 'redis' at host:port 172.16.22.21:12000.
- Data streams overview:** Provides an overview of data streams. For the 'Total' stream, there are 350 total items, 0 pending, 95 inserted, 0 updated, 0 deleted, 255 filtered, 0 rejected, 0 deduplicated, and the last arrival was 06 Dec 2024 at 17:27:51.437. The 'rdiPublicTrack' stream has similar statistics.
- Clients:** Lists connected clients. There are two clients: one with ID 512379001002 at address 172.16.22.1:37065 and another with ID 512386001001 at address 172.16.22.1:47416, both with N/A names and users.

# Processing performance

Includes these statistics

The screenshot shows the Redis Dashboard interface. At the top, there are navigation icons for Home, Overview, Configuration, and Help. Below the header, the RDI Version is listed as 7.2.4, the Address is 172.16.22.21:12001, the Run status is ready, and the Sync Mode is streaming. On the left, there's a sidebar with a Processing performance information section, which is currently expanded. This section contains six performance metrics:

Metric	Value	Unit
Total batches	274	Total
ACK time average	0.4	sec
Batch size average	1.7	MB
Records per second average	16.1	/sec
Process time average	4.9	ms
Read time average	100.2	ms

To the right of these metrics, there is a dropdown menu with the text "Auto refresh: 5.0 s" followed by a circular refresh icon and a dropdown arrow. A yellow box highlights this area.

At the bottom of the dashboard, there is a "Total time average" card. Inside this card, there is a "Auto Refresh" toggle switch, which is turned on, and a "Refresh rate: 5.0 s" setting. A yellow box highlights this entire card, and a yellow arrow points from the top-right dropdown to this card, with the text "Can be set to static or auto refresh" written next to the arrow.

# Target connections

Connected

Target connections			
Status	Name ↑	Type	Host:port
●	target	redis	172.16.22.21:12000

Disconnected

Target connections			
Status	Name ↑	Type	
⚠	targetDb	redis	

# Data streams overview

Data streams overview									
Stream name ↑	Total	Pending	Inserted	Updated	Deleted	Filtered	Rejected	Deduplicated	Last arrival
Total	2950	2	732	0	0	2218	0	0	
rdiPublicTrack	2950	2	732	0	0	2218	0	0	06 Dec 2024 17:41:02.838

# Collector metrics

The endpoint for collector metrics is [https://<RDI\\_HOST>/metrics/collector-source](https://<RDI_HOST>/metrics/collector-source)

These metrics are divided into three groups:

- Pipeline state: metrics about the pipeline mode and connectivity
- Data flow counters: counters for data breakdown per source table
- Processing performance: processing speed of RDI micro batches

# Stream processor metrics

The endpoint for stream processor metrics is [https://<RDI\\_HOST>/metrics/rdi](https://<RDI_HOST>/metrics/rdi)

RDI reports metrics during the two main phases of the ingest pipeline, the snapshot phase and the change data capture (CDC) phase.

[The full list is in the docs.](#)

# Grafana dashboard for RDI



# Grafana dashboard for RDI

## Job data

- Pending - how many records are waiting to be written to Redis?
- Updated - Number of updated key values (i.e. completed update operations)
- Deleted - Number of deleted keys
- Inserted - Number of inserted keys
- Filtered - Number of key changes skipped because they were filtered out
- Rejected - Number of errors - ideally, always zero in production

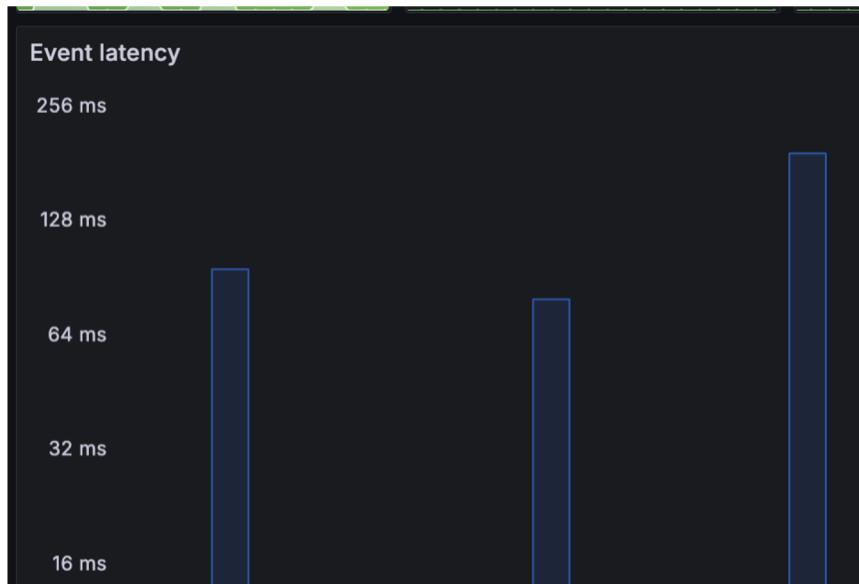
### Job data for rdi.public.Track



# Grafana dashboard for RDI

You can also see latency and ops/sec

Event latency



Operations per second



sd2151kd1fq05m18iq10zmahtrw2mvmlbasj35jk112pldfu  
n457n39dife29f9x5v5g66rga7sew9dcg6yji3lp2a4qwe8r  
t9yuip6dfg4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215  
ikd17**thankyou**.rw2mvmlbasj35jk112pldfun457n39dif  
e29f9x5v5g66rga7sew9dcg6yji3lp2a4qwe8rt9yuip6dfg  
4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m1  
8iq10zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9  
x5v5g66rga7sew9dcg6yji3lp2a4qwe8rt9yuip6dfg4hjk4  
11mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m18iq10  
zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g  
66rga7sew9dcg6yji3lp2a4qwe8rt9yuip6dfg4hjk411mnb  
2vcx **Redis** d5f68fet4y5h6sd215ikd1fg05m18iq10zmdhf  
rw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g66rga  
7sew9dcg6yji3lp2a4qwe8rt9yuip6dfg4hjk411mnb2vcx3

# General troubleshooting



# General troubleshooting

- `redis-di status` command - discover source, target, data flow (next slide)
- Logs (following slide)
- Debezium metrics - see the state of collection
- Support package - has everything
- If you want to run kubectl to get immediate components' status: use `-n rdi` with all Kubectl commands
  - Start with  
`kubectl -n rdi get all -a`
  - Then you can continue with  
`kubectl -n rdi describe [pod|configmap|secret]<object name>`

# redis-di status args [OPTIONS]

These are the arguments and options

- args
  - rdi\_host (REQUIRED)
  - rdi\_port (REQUIRED)
  - rdi\_user (if applicable)
  - rdi\_password (if applicable)
- Options
  - log\_level [DEBUG, INFO, WARN, ERROR, CRITICAL]
    - redis-di status --log-level DEBUG
  - live (for live data flow monitoring)
    - redis-di status --live TRUE
  - redis-di status --help

# Logs

RDI uses fluentd and logrotator to ship and rotate logs at /opt/rdi/logs

Component	Log name
Operator	/opt/rdi/logs/operator.log
Stream Processor	/opt/rdi/logs/processor.log
Other logs	/opt/rdi/logs/monitor.log

# Logs

“Other” logs, continued

Filename	Phase
rdi_collector-collector-initializer.log	Initializing the collector
rdi_collector-debezium-ssl-init.log	Establishing the SSL connections to the source and RDI database (if you are using SSL)
rdi_collector-collector-source.log	Collector change data capture (CDC) operations
rdi_rdi-rdi-operator.log	Main RDI control plane component
rdi_processor-processor.log	RDI stream processing

# Metrics

## Getting access to metrics

- Debezium collects metrics and we expose them through a dedicated Prometheus exporter
- The Stream Processor writes metrics to RDI database and the Metrics exporter provides these metrics
- The CLI status command can provide live or snapshot metrics
- There is an API to get metrics as well

## Which metrics are available:

- Debezium has different metrics per phase of the pipeline:
  - Snapshot metrics
  - CDC/Streaming metrics
- The Stream Processor reports data counters per source table and performance metrics per micro batch

sd2151kd1fq05m18iq10zmahtrw2mvmlbasj35jk112pldfu  
n457n39dife29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8r  
t9yuip6dfg4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215  
ikd17**thankyou**.rw2mvmlbasj35jk112pldfun457n39dif  
e29f9x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg  
4hjk411mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m1  
8iq10zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9  
x5v5g66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk4  
11mnb2vcx3za4s5d5f68fet4y5h6sd215ikd1fg05m18iq10  
zmdhfrw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g  
66rga7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb  
2vcx **Redis** d5f68fet4y5h6sd215ikd1fg05m18iq10zmdhf  
rw2mvmlbasj35jk112pldfun457n39dife29f9x5v5g66rga  
7sew9dcg6yji31p2a4qwe8rt9yuip6dfg4hjk411mnb2vcx3